

PRD: Eliq Data Platform

1. Objective

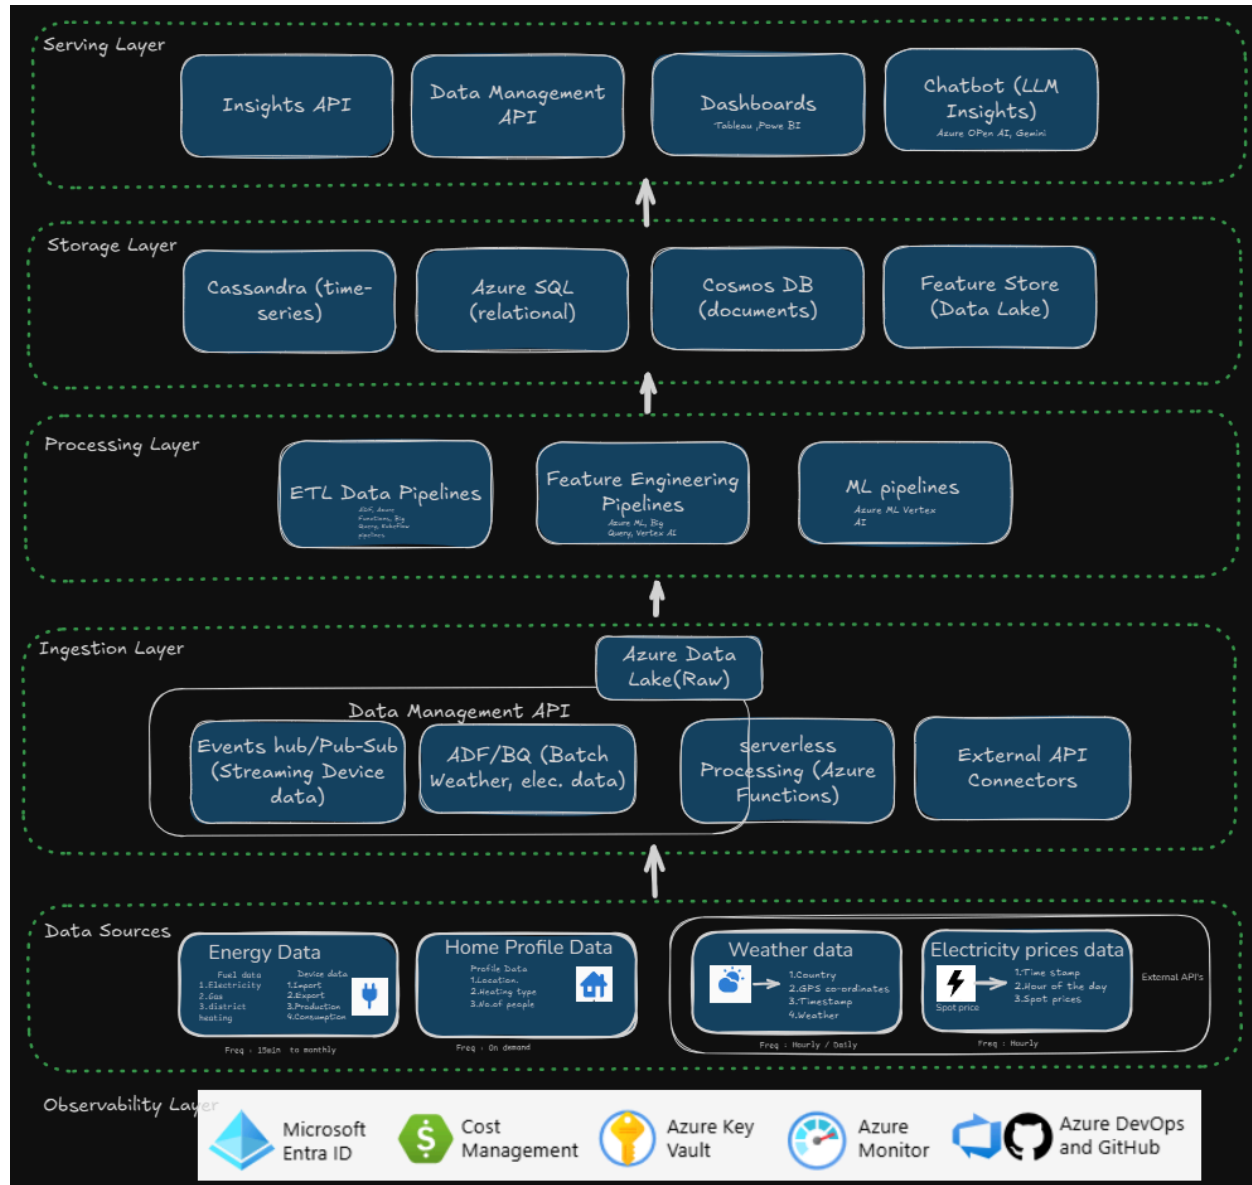
Build a **scalable, cloud-native data platform** that:

- Ingests and integrates **energy, home, weather, and market data**.
 - Provides a **single source of truth** for analytics and APIs.
 - Enables internal teams to **deliver insights quickly and reliably**.
-

2. Key Goals

- **Unified ingestion** from heterogeneous sources.
- **Data quality & governance** at every stage.
- **Faster insights** for internal teams and clients.
- **Modular architecture** that can scale as data grows.

Data Architecture Diagram



Component-by-Component Explanation

1. Data Sources

- **Energy Data:** From devices like meters or inverters, capturing import/export/production/consumption. Granularity varies from monthly to 15-min intervals.
 - **Home Profile Data:** Stored as structured documents in Cosmos DB, describing physical/occupancy attributes per household.
 - **Weather Data:** Hourly data for fixed GPS coordinates via 3rd-party weather API.
 - **Electricity Price Data:** Spot prices pulled daily via API for hourly intervals.
 - **Internal Insight Services:** Compute on-the-fly results without persisting—used as ephemeral data processors.
-

2. Ingestion Layer

Data Management API handles the data ingestion process from various data sources while External API connectors helps to connect to extract the Weather and elec. data.

- **Azure Data Factory/Big Query:** Handles scheduled pulls for weather and price data.
 - **Event Hub/Pub Sub + Azure Functions:** Manages real-time or streaming data from energy devices.
 - **Schema validation & transformation:** Light checks and standardization on ingestion.
-

3. Raw Data Lake (Azure Data Lake Gen2/Big Query)

- Central landing zone.
- Immutable, append-only storage.

- Organized by source, date, and location.
-

4. Processing Layer (Azure Functions,Big Query,Vertex AI,Azure ML)

- ETL Pipelines - Cleans, joins, enriches raw data:
 - Combines energy with weather and home profiles.
 - Aggregates to common time resolutions.
 - Feature Engineering Pipelines - Prepares features for modeling.
 - Hosts ML models: forecasting, disaggregation, similarity scoring.
-

5. Curated Storage

- **Cassandra**: High-write, distributed time-series store for granular energy data.
 - **Azure SQL**: Relational queries for reports and summaries.
 - **Cosmos DB**: Retains document-format home profiles.
 - **Feature Store**: Centralized, versioned features for model reproducibility.
-

6. Serving Layer

- **Insights API**: Real-time access to computed insights per location.
 - **Data Management API**: Secure, structured ingestion endpoint for partners.
 - **Tableau/Power BI Dashboards**: Internal tools for Customer Success, Sales, and Ops.
 - **Chatbot/LLM based Insights : Enable** users to leverage automated insights using LLMs.
-

Integration Approach with Data sources

Source Type	Integration Method	Frequency
Energy data	Event Hub or secure SFTP + Data Factory	Real-time / Daily
Home profile data	Cosmos/Big Query DB (direct sync/API)	On-demand
Weather data	REST API via Azure functions/Data Factory	Hourly / Daily
Market prices	API pull via Functions/Data Factory	Daily
Internal services	Invoked at API request (stateless)	Real-time

Responsibilities by Team ,Layer,Tools

- **Data Engineering (2-3 people):** Drives ingestion pipelines, data storage.
- **Data Science (1-2 people):** Drives ML and feature creation.
- **Platform Engineering(1-2 people):** Owns APIs and real-time services.
- **Cloud/DevOps(1 people):** Ensures infrastructure, monitoring, CI/CD.
- **Customer Success(1 people):** Focused on dashboards and insights delivery.
- **Security/Compliance(1 people):** Governs access, policies, audits.

RACI Matrix

Key:

- **R** = Responsible (does the work)
- **A** = Accountable (owns the outcome)
- **C** = Consulted (provides input)
- **I** = Informed (kept up to date)

Component / Layer	Data Eng	Data Science	Platform Eng	Cloud / DevOps	Customer Success	Security / Compliance
Data Sources	I	I	I	I	I	I
Event Hub / Streaming	R/A	C	C	C	I	I
ADF / Batch Pipelines	R/A	C	I	C	I	I
Serverless Processing (Azure Functions)	R	C	A	C	I	I
External API Connectors	R	C	A	C	I	I
Data Management API	C	I	R/A	C	I	I
Raw Azure Data Lake	R	C	I	A	I	C
ETL Data Pipelines	R/A	C	I	C	I	I
Feature Engineering Pipelines	C	R/A	I	C	I	I
ML Pipelines	C	R/A	I	C	I	I
Cassandra / Azure SQL / Cosmos DB	R/A	C	I	C	I	C
Feature Store	C	R/A	I	C	I	C
Insights API	C	C	R/A	C	I	I
Dashboards (Power BI / Tableau)	C	C	I	I	R/A	I
Chatbot (LLM Insights)	C	R/A	C	I	I	I
Identity & Access (Microsoft Entra ID)	I	I	I	R/A	I	C
Key Vault / Security Config	I	I	I	R/A	I	C
Cost Management / Azure Monitor	I	I	I	R/A	I	C

How Teams Interact with the Data Platform

Team	How They Use the Platform	Key Components They Rely On
Data Science	<ul style="list-style-type: none">- Consume curated data for analysis and modeling- Build ML pipelines (forecasting, disaggregation, similar homes)- Deploy models exposed via Insights API	Feature Store, Databricks, ML Pipelines, Cosmos DB
Data Engineering	<ul style="list-style-type: none">- Build ingestion pipelines (Event Hub, Data Factory)- Maintain ETL workflows- Manage curated storage and schemas	Event Hub, Data Factory, Azure Data Lake, ETL Pipelines, Cassandra, Azure SQL, Cosmos DB
Platform / Backend Eng.	<ul style="list-style-type: none">- Develop and manage APIs (Data Management API & Insights API)- Ensure platform uptime and integrations	Data Management API, Insights API, Azure Functions
Customer Success / Analytics	<ul style="list-style-type: none">- Use dashboards to view and explain insights to clients- Provide feedback on model accuracy and data quality	Dashboards (Power BI/Tableau), Insights API
Sales / Business Dev.	<ul style="list-style-type: none">- Use dashboards and APIs for client demos- Present data- driven insights to prospects	Dashboards, Insights API

Approach if Plan Approved

1. Kick-off sprint with Data Engineering to set up cloud infrastructure.
2. Deliver incremental milestones every 2 weeks (Agile).
3. Conduct reviews after each phase with relevant stakeholders (Engineering, Data Science, Customer Success).
4. Final sign-off once APIs and dashboards meet defined acceptance criteria.

Total Team Size (MVP phase):

- Involving 7–10 people is sufficient for 8–10 weeks of MVP delivery.

Initial User Stories and Tasks

User Story	Tasks	Acceptance Criteria	Review & Sign-off
As a Data Engineer, I can ingest 15-minute energy data into the data lake	- Create Event Hub - Build Functions for validation - Store raw data in Data Lake	- Energy data arrives in <i>/raw/energy/</i> partitioned folders - Validation errors logged properly	Reviewed by Data Eng Lead; sign-off by Platform Architect
As a Data Scientist, I can access curated datasets that combine energy, weather & profiles	- Build batch pipelines (ADF/Databricks) - Join energy, weather & home profiles - Create curated tables	- Curated table available in Data Lake and SQL - Data refresh automated daily	Reviewed by Data Science Lead
As a Developer, I can query aggregated energy data via Insights API	- Expose curated data to API - Build endpoint for aggregated consumption queries	- API returns correct aggregation results within SLA	Reviewed by Platform Eng Lead
As a Customer Success Manager, I can view consumption trends on a dashboard	- Connect Power BI to curated datasets - Create basic consumption trends dashboard	- Dashboard shows correct time-series for selected locations	Reviewed by Customer Success Lead

Execution Plan

Phase & Timeline (approx.)	Key Activities	Outcome Deliverables
Week 1–2: Foundation	<ul style="list-style-type: none">- Provision Azure Data Lake- Set up Event Hub & Data Factory- Define raw data schemas	Raw landing zone & ingestion setup
Week 3–4: Ingestion	<ul style="list-style-type: none">- Build pipelines for energy, weather, market data- Implement validation & logging	Automated data ingestion running
Week 5–6: Processing	<ul style="list-style-type: none">- Develop Databricks ETL pipelines- Curate datasets (energy + weather + profiles)	Curated tables ready for consumption
Week 7–8: Serving	<ul style="list-style-type: none">- Connect curated data to Insights API- Create dashboards in Power BI	APIs live and dashboards available
Ongoing	<ul style="list-style-type: none">- Add feature store- Build ML pipelines- Enhance dashboards	Iterative improvements (V1 → V2)